

# ECE 592 Project 1

Kudiyar Orazymbetov (korazym@ncsu.edu)  
Nico Casale (ncasale@ncsu.edu)

September 25, 2017

Note that the entries are links.

## Contents

<b>1</b>	<b>Introduction</b>	<b>3</b>
<b>2</b>	<b>Loading an image</b>	<b>3</b>
<b>3</b>	<b>Reconstruction of an Image with Assigned Clusters</b>	<b>4</b>
<b>4</b>	<b>Rate vs. Distortion Performance</b>	<b>4</b>
<b>5</b>	<b>Varied Patch Sizes</b>	<b>5</b>
<b>6</b>	<b>Better Compression with Entropy Coding</b>	<b>6</b>
<b>7</b>	<b>Additional Direction: MATLAB implementation of K-Means</b>	<b>6</b>
7.1	Initialization of <i>k-means</i> . . . . .	7
7.2	Timing and Distortion Comparisons . . . . .	7
7.3	Profiler Results . . . . .	8
<b>8</b>	<b>Conclusion</b>	<b>9</b>
<b>9</b>	<b>Appendix A: Code Listings</b>	<b>10</b>
9.1	Code for Parts 2-5 of the Project . . . . .	10
9.1.1	<code>mainCode.m</code> . . . . .	10
9.2	Code for our implementation of K-means . . . . .	12
9.2.1	<code>compare_Kmeans.m</code> . . . . .	12
9.2.2	<code>kmeans_alt.m</code> . . . . .	14
9.2.3	<code>assignmentStep.m</code> . . . . .	15
9.2.4	<code>updateStep.m</code> . . . . .	16

## List of Tables

1	Coding Rate vs. Normalized Coding Length with Entropy Coding. . . . .	6
2	Comparison of MATLAB's <i>k-means</i> and ours. . . . .	7

## List of Figures

2.1	One of the images we used for the project. . . . .	3
3.1	The original image (left) with a quantized version (right). . . . .	4
4.1	Rate (R) and Distortion (D) performance for patch size $P = 2$ . . . . .	5
5.1	Rate (R) and Distortion (D) performance for patch size $P = 2, 4$ . . . . .	5
7.1	Example output of <code>kmeans_alt(.)</code> . . . . .	7
7.2	First page of MATLAB's profiler results. . . . .	8
7.3	Analysis of the <code>assignmentStep(.)</code> function. . . . .	8

## Listings

1	Code to solve parts 2-5 of the Project. . . . .	10
2	Code to compare our implementations of <i>k-means</i> . . . . .	12
3	Our implementation of <i>k-means</i> . . . . .	15
4	Helper function for <code>kmeans_alt(.)</code> . . . . .	15
5	Helper function for <code>kmeans_alt(.)</code> . . . . .	16

## 1 Introduction

This project covers an implementation and study of image compression using *k-means* clustering. We show that an image can be reconstructed from representative patches that are learned from its particular distribution of pixels. Qualitatively, the reconstruction is not particularly unappealing and the quantitative results are not poor, either.

As a further experiment, we generated our own implementation of *k-means* in MATLAB that is qualitatively similar to the output of MATLAB's built in `kmeans(.)` function but does not yet compete in terms of time to execute and quantitative distortion values.

The code used to execute this project can be found in Appendix A.

## 2 Loading an image

For the project we chose an image of size 512x512, pictured below. Note that we used a grayscale image for a low-complexity implementation that computes quickly. If we used a color image, the *k-means* algorithm would be passed a matrix that was 3× as large, causing it to converge slowly.



Figure 2.1: One of the images we used for the project.

### 3 Reconstruction of an Image with Assigned Clusters



Figure 3.1: The original image (left) with a quantized version (right).

Qualitatively, the image is recognizable and not dramatically distorted by the reconstruction. This image is the result of  $P = 2$  and  $K = 16$ .

### 4 Rate vs. Distortion Performance

For this part, we chose rate  $R$  to be  $[0.25, 0.5, 0.75, 1]$ . To find the number of clusters given a certain rate, we used the following:

$$C = \text{round}(2^{RP^2}) \quad (4.1)$$

Where  $C$  is the number of clusters,  $R$  is the coding rate, and  $P$  is the patch size (in one dimension only.)

The corresponding number of clusters to the rates given is  $[2, 4, 8, 16]$  for  $P = 2$ . The figure below illustrates the results of these variations.




Figure 4.1: Rate (R) and Distortion (D) performance for patch size  $P = 2$

The plot shows that as we increase the rate, distortion decreases. This corresponds to an increase in cluster numbers, which allows the *k-means* algorithm to more accurately represent the subtleties of the image.

## 5 Varied Patch Sizes

As our image size is 512x512, the maximum number of clusters, which is power of 2, for the case of  $P = 4$  is 16384. If we used 16384 clusters, each patch of the image could be represented by its unique patch, which corresponds to the rate of 0.75. Because we cannot represent the image in more clusters than actually exist in the image, we only attempted rates from [0.25, 0.5, 0.75]. Below is a graph of the original RD plot with the new RD values for  $P = 4$ .

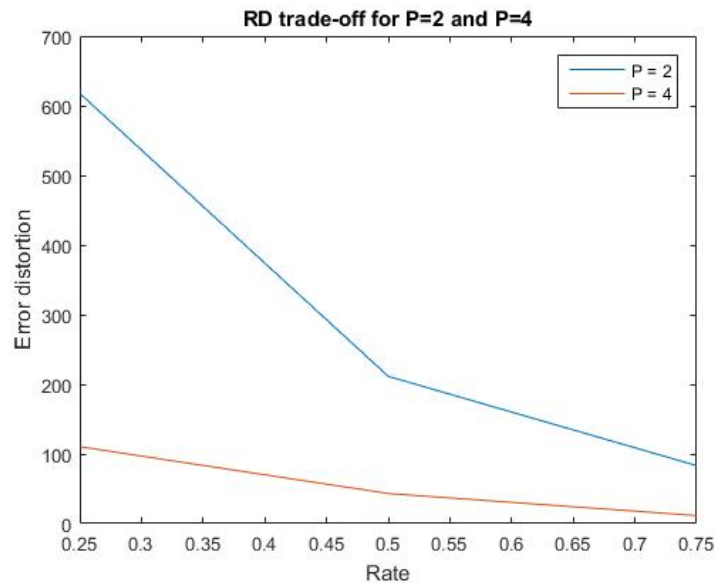


Figure 5.1: Rate (R) and Distortion (D) performance for patch size  $P = 2, 4$

From the plot, we can see that even with a small rate increase we can substantially decrease the error distortion. It has a steeper drop in error when we increase the rate from 0.25 to 0.5 compared to the increase from 0.5 to 0.75. The underlying reason of getting less error with a patch size of 4 is we are having 16 components for each patch. Thus, the *k-means* algorithm has more dimensions to calculate the distance to other patches. Each cluster can convey a wider range of values and capture fine-grained detail in pixel change across patch boundaries. When  $P = 2$  there are only 4 pixels in each patch, so the expressive capabilities of *k-means* clustering is limited.

## 6 Better Compression with Entropy Coding

Being that some clusters are used more frequently to represent patches in the image, we can reduce the coding length for more likely clusters. This is similar to Morse code, where the most common letters are represented in the smallest numbers of dots and dashes. For different coding rates  $R$ , which correspond to the number of clusters as Eq. (4.1), we noticed that we could reduce the number of bits required to represent each cluster (in a normalized measure) in each case.

<b>R</b> <b>(Coding Rate)</b>	<b>Normalized Coding Length</b> <b>(bits per pixel)</b>
0.25	0.250
0.5	0.482
0.75	0.735
1	0.962

Table 1: Coding Rate vs. Normalized Coding Length with Entropy Coding.

From above we can conclude the bits we use per pixel is smallest when we have a smaller cluster size ( $C = 2$ ). As we increase the cluster size, we need to use more bits to represent each cluster. Nevertheless, we can decrease the bits for 1 to 0.96 when cluster size is 16. In this configuration, we can achieve a 4% reduction in coding rate over an implementation that does not use *entropy coding*.

## 7 Additional Direction: MATLAB implementation of K-Means


To develop an implementation of *k-means* in MATLAB, we referenced Wikipedia's [article on K-means](#).

The *k-means* algorithm (*a.k.a.* Lloyd's algorithm) has three main steps:

1. Initialize the  $k$  centroids, or cluster locations.
2. Assign each sample of the data to one of the centroids. The centroid with the smallest squared Euclidian distance to the sample is chosen.
3. Update each cluster as the centroid (geometric average in  $P^2$ -Dimensional space, where  $P$  is the patch dimension) of the samples that were assigned to it in step 2.
4. repeat steps 2 and 3 until the assignments no longer change.

Note that *k-means* is not guaranteed to find the global optimum clustering, as it is highly dependent on step 1. The initialization is the primary factor in the convergence and quality of the clustering operation. Some notes on initialization are in the next section.

Below is an example of the output of our *k-means* function on a  $[400 \times 400]$  image with a patch size of  $[2 \times 2]$  and 16 clusters.

Figure 7.1: Example output of `kmeans_alt(.)`.

## 7.1 Initialization of *k-means*

At first, we took the following approach to initialize the clusters:

1. Find the global minimum and maximum of the pixels in the image.
2. Generate  $K$  linearly spaced points in  $\mathbb{R}^{1 \times P}$ .

However, this approach took  $\sim 50$  iterations to converge. Likewise, because it didn't take into account the statistical distribution of the pixels in the image, it couldn't capture the subtleties of the image where most of its pixels lay.

After looking at the Wikipedia page for *k-means*, we noticed that there was some discussion on initialization methods. As a second approach, we tried the Forgy Method, which takes  $k$  random samples from the image and uses them as the initial clusters. This method was much more effective as it reduced the number of iterations needed to converge to  $\sim 15$ . As this method is inherently stochastic, some assays would take longer to converge as they were initialized to locations that were distant from the more stable regions. We describe good locations for clusters as 'stable regions' because there are an infinite number of clusters and locations that could converge under Lloyd's algorithm.

## 7.2 Timing and Distortion Comparisons

The table below illustrates some comparisons between MATLAB's built in *k-means* function and the one we coded. All values were averaged over 20 iterations in MATLAB R2016b. The computer used was a Dell laptop with an Intel Core i3 processor.

	MATLAB	Our Implementation
Average time	0.863	2.578
Minimum time	0.421	0.965
Distortion	27.386	56.191

Table 2: Comparison of MATLAB's *k-means* and ours.

Note that our function is both slower and less accurate than MATLAB's by a factor of  $\sim 2$ . In the future, improvements could be made to the initialization procedure to try to get a more accurate representation that takes less time to compute. From a purely computational perspective, the *k-means* exhibits a great deal of parallelism. The result of step 3 is strictly dependent on step 2, but each is parallelizable on its own. In the next section, some details on a potential parallelization of *k-means* are discussed.

### 7.3 Profiler Results

In analyzing the bottlenecks of the code we wrote to implement *k-means*, we noticed that there is one computationally expensive part of the calculation that is hard to improve outside of parallelization. The figures below illustrate the most time-consuming lines of our function. Note that this screenshot is from a run of our alternate *k-means* function that converged in 9 iterations, a relative minimum given our current implementation.




Figure 7.2: First page of MATLAB's profiler results.

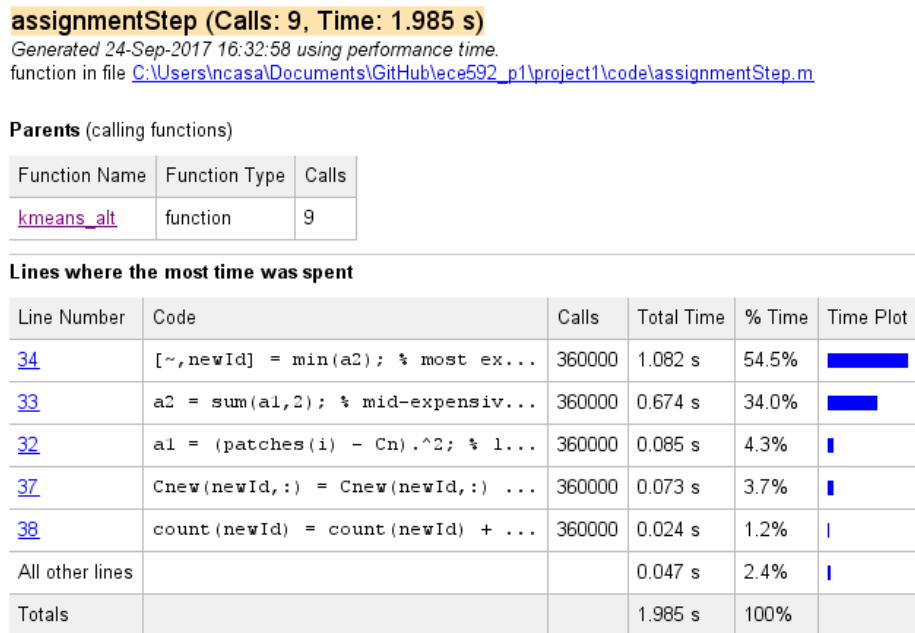


Figure 7.3: Analysis of the `assignmentStep(.)` function.

As shown in the figure, our two most expensive lines deal with the assignment of samples to clusters. In order to determine which of the current clusters a sample belongs to, the code takes the squared Euclidian distance of each point to each of the clusters. The cluster that is closest to the point is assigned. This step is computationally expensive because the function `assignmentStep(.)` must sum along the y dimension the distances between each of the pixel values in the patch and the centroids. Then, the `min(.)` function has to find the smallest point in the



$[K \times 1]$  array. These calls are necessary for quick convergence but negatively impact the time it takes to reach convergence. We tried to use ersatz methods, such as taking the absolute value of the distances instead of squaring them, or taking only the 1<sup>st</sup> column of distances, but these approaches led to less accurate assignment. Even though they were less computationally expensive individually, they led to a higher number of iterations, which totally counteracted the positive effects of the individual computation.

If we were to implement *k-means* on a GPU, we could parallelize the thousands of calls to `assignmentStep(.)` so that the function would run faster. That, paired with a more effective initialization, would create a good implementation of *k-means* that could be on par with MATLAB.

## 8 Conclusion

In conclusion, we have found that it is possible to compress an image using clustering algorithms such as *k-means*. The quality of the reconstruction is dependent on the patch size  $P$  and the number of clusters  $C$  ( $K$  in our MATLAB code.) There is also a reasonable level of additional compression that could be added if we used entropy coding to represent more common clusters with less bits. Finally, we have seen that a capable implementation of *k-means* is possible and not especially complex.

## 9 Appendix A: Code Listings

### 9.1 Code for Parts 2-5 of the Project

#### 9.1.1 mainCode.m

Listing 1: Code to solve parts 2-5 of the Project.

```

1  %{
2  ECE 592 Project 1
3
4  Kudiyar (Cody) Orazymbetov
5  korazym@ncsu.edu
6
7  Nico Casale
8  ncasale@ncsu.edu
9  %}
10
11 %% PART 1
12 clear; close all;
13 addpath('utility', 'kmeans');
14 setup();
15 global imagesFolder
16 overwriteImage = 0;
17
18 fprintf('ECE 592 Project 1\n');
19 fprintf(strcat(datestr(now),'\n'));
20 % read image
21 I = double(imread('image3.gif'));
22 [M, N] = size(I);
23
24 %% PART 2
25 P = 2; % square patch dimension
26 R = 0.25:0.25:1; % various values for Rate
27 K = round(2 .^(R*P^2));
28 for j = 1:length(K)
29     % partition into patches
30     Ipartitioned = im2col(I, [P P], 'distinct');
31
32     % apply k-means
33     [idx, Cn] = kmeans(Ipartitioned', K(j));
34
35     % reconstruct image
36     indexrepresentations = zeros(length(idx), P^2);
37     for i = 1:length(idx)
38         indexrepresentations(i, :) = Cn(idx(i), :);
39     end
40     Iquantized = col2im(indexrepresentations', [P P], size(I), 'distinct');
41
42     subplot(1, 2, 1);
43     imshow(I, []);
44     f2 = subplot(1,2,2);
45     imshow(Iquantized, []);
46
47     % Calculate distortion D
48     D(j) = sum(sum((I - Iquantized) .^2)/(M*N));
49 end

```

```

50
51 %% PART 3 Rate & Distortion Plot
52 R = 0.25:0.25:1;
53 plot(R,D)
54
55 %% PART 4
56 P = 4;
57 R = 0.25:0.25:0.75;
58 K = round(2 .^(R*P^2));
59 for j = 1:length(K)
60     % partition into patches
61     Ipartitioned = im2col(I, [P P], 'distinct');
62
63     % apply k-means
64     [idx, Cn] = kmeans(Ipartitioned', K(j));
65
66     % reconstruct image
67     indexrepresentations = zeros(length(idx), P^2);
68     for i = 1:length(idx)
69         indexrepresentations(i, :) = Cn(idx(i), :);
70     end
71     Iquantized = col2im(indexrepresentations', [P P], size(I), 'distinct');
72
73     % compute distortion D
74     D1(j) = sum(sum((I - Iquantized) .^2)/(M*N));
75 end
76
77 % Plot RD trade-off for P =2 and P =4;
78 figure
79 R = 0.25:0.25:0.75;
80 plot(R(1:3),D(1:3))
81 title('RD trade-off for P=2 and P=4')
82 xlabel('Rate')
83 ylabel('Error distortion')
84 hold on
85 plot(R,D1)
86 hold off
87 legend('P = 2', 'P = 4')
88
89 %% PART 5
90 % First run Part 2, then run this part
91
92 P = 2; % square patch dimension
93 R = 0.25:0.25:1;
94 K = round(2 .^(R*P^2));
95
96 for j = 1:length(K)
97     % partition into patches
98     Ipartitioned = im2col(I, [P P], 'distinct');
99
100     % apply k-means
101     [idx, Cn] = kmeans(Ipartitioned', K(j));
102
103     % recontstruct image
104     indexrepresentations = zeros(length(idx), P^2);
105     for i = 1:length(idx)

```

```

106     indexrepresentations(i, :) = Cn(idx(i), :);
107 end
108 Iquantized = col2im(indexrepresentations', [P P], size(I), 'distinct');
109
110 % plot results
111 subplot(1, 2, 1);
112 imshow(I, []);
113 f2 = subplot(1,2,2);
114 imshow(Iquantized, []);
115
116 % compute distortion
117 D(j) = sum(sum((I - Iquantized) .^2)/(M*N));
118
119 for k = 1:K(j)
120     Np(k) = sum(idx == k);
121 end
122
123 H = 0;
124 for m = 1:K(j)
125     H = H - (Np(m)*P^2/(M*N)*log2(Np(m)*P^2/(M*N)));
126 end
127
128 r(j) = H/P^2; % look at r array to see rates
129 end

```

## 9.2 Code for our implementation of K-means

### 9.2.1 compare\_Kmeans.m

Listing 2: Code to compare our implementations of *k-means*.

```

1  %{
2  ECE 592 Project 1
3
4  Kudiyar (Cody) Orazymbetov
5  korazym@ncsu.edu
6
7  Nico Casale
8  ncasale@ncsu.edu
9
10 This file is used to compare the built-in kmeans function and
11 the function generated by Cody and Nico for project 1.
12 %}
13
14
15 clear; close all;
16 addpath('../utility', '..');
17 global seed
18 seed = 475859;
19 rng(seed);
20 global imagesFolder
21 imagesFolder = '../images';
22 addpath(imagesFolder);
23 overwriteImage = 0;
24
25 fprintf('ECE 592 Project 1\n');

```

```

26 fprintf(strcat(datestr(now),'\n'));
27
28 % read image
29 sz = 400;
30 I = double(rgb2gray(imcrop(imread('cassava.jpg'),[50 50 sz-1 sz-1])));
31 [M, N] = size(I);
32 %image(I) properties
33 %whos I
34 f = figure(1);
35 I3 = cat(3, I, I, I);
36 image(uint8(I3));
37 prettyAxes();
38 prettyPictureFig(f); prettyPictureFig(f);
39
40 P = 2; % square patch dimension
41
42 % partition into patches
43 Ipartitioned = im2col(I, [P P], 'distinct');
44
45 %% MATLAB's kmeans
46 % apply k-nearest neighbor
47 K = 16;
48 tic;
49 [idx, Cn] = kmeans(Ipartitioned', K);
50 toc;
51
52 % recreate image with the clusters learned by kmeans
53 indexrepresentations = zeros(length(idx), P^2);
54 for i = 1:length(idx)
55     indexrepresentations(i, :) = Cn(idx(i), :);
56 end
57
58 Iquantized = col2im(indexrepresentations', [P P], size(I), 'distinct');
59
60 f = figure(2);
61 quantizedI3 = cat(3, Iquantized, Iquantized, Iquantized);
62 image(uint8(quantizedI3));
63 prettyAxes();
64 prettyPictureFig(f);prettyPictureFig(f);
65
66
67 %% Nico and Cody's kmeans
68 profile on
69 tic
70 [idx_alt, Cn_alt] = kmeans_alt(Ipartitioned', K);
71 toc
72 profile viewer
73
74 indexrepresentations_alt = zeros(length(idx_alt), P^2);
75 for i = 1:length(idx_alt)
76     indexrepresentations_alt(i, :) = Cn_alt(idx_alt(i), :);
77 end
78
79 Iquantized_alt = col2im(indexrepresentations_alt', [P P], size(I), 'distinct');
80
81 f = figure(3);

```

```

82 quantizedI3_alt = cat(3, Iquantized_alt, Iquantized_alt, Iquantized_alt);
83 image(uint8(quantizedI3_alt));
84 prettyAxes();
85 prettyPictureFig(f);prettyPictureFig(f);
86
87 compare = [Cn Cn_alt];
88
89 D_kmeans = sum(sum((I(:,:,1) - quantizedI3(:,:,1)) .^2)/(M*N))
90 D_kmeans_alt = sum(sum((I(:,:,1) - quantizedI3_alt(:,:,1)) .^2)/(M*N))
91
92 %% Get an average time and distortion for each
93 REPS = 20;
94 tMin = inf;
95 tic;
96 for rep = 1:REPS
97     fprintf('rep = %d\n', rep);
98     rng shuffle
99     tSt = tic;
100     [idx, Cn] = kmeans(Ipartitioned', K);
101     tElapsed = toc(tSt);
102     tMin = min(tElapsed, tMin);
103
104     D_kmeans(rep) = sum(sum((I(:,:,1) - quantizedI3(:,:,1)) .^2)/(M*N));
105
106 end
107 tAvg = toc/REPS;
108 fprintf('MATLAB: min = %d, avg = %d\n', tMin, tAvg);
109
110 REPS = 20;
111 tMin_alt = inf;
112 tic;
113 for rep = 1:REPS
114     fprintf('rep = %d\n', rep);
115     rng shuffle
116     tSt = tic;
117     [idx_alt, Cn_alt] = kmeans_alt(Ipartitioned', K);
118     tElapsed_alt = toc(tSt);
119     tMin_alt = min(tElapsed_alt, tMin_alt);
120
121     D_kmeans_alt(rep) = sum(sum((I(:,:,1) - quantizedI3_alt(:,:,1)) .^2)/(M*N));
122
123 end
124 tAvg_alt = toc/REPS;
125 fprintf('592: min = %d, avg = %d\n', tMin_alt, tAvg_alt);
126
127 D_kmeans_avg = mean(D_kmeans);
128 D_kmeans_avg_alt = mean(D_kmeans_alt);
129 fprintf('Distortion - MATLAB = %d, 592 = %d\n', D_kmeans_avg, D_kmeans_avg_alt);
130
131 %% save images
132 file = sprintf('altkmeans_reconstructed');
133 file = strcat(imagesFolder, file);
134 print(file, '-dpng');

```

### 9.2.2 kmeans\_alt.m

Listing 3: Our implementation of *k-means*.

```

1  %{
2  ECE 592 Project 1
3
4  Kudiyar (Cody) Orazymbetov
5  korazym@ncsu.edu
6
7  Nico Casale
8  ncasale@ncsu.edu
9
10 This is our own implementation of K-means
11 https://en.wikipedia.org/wiki/K-means_clustering#Standard_algorithm
12 %}
13
14 function [idx, Cn] = kmeans_alt(patches, K)
15
16     %{
17     Generate initial centroids by choosing
18     K random samples of patches
19     %}
20     inds = randi(length(patches), [K, 1]);
21     Cn = patches(inds,:);
22
23     % Group and update centroids
24     iter = 0;
25     max_iter = 100; % maximum iterations
26
27     change = inf;
28     idx = zeros(length(patches),1);
29
30     while (change && iter < max_iter)
31         [Cnew, count, idx] = assignmentStep(Cn, idx, patches);
32         Cnew = updateStep(Cnew, count);
33         change = max(max(abs(Cnew - Cn)));
34         Cn = Cnew;
35         iter = iter + 1;
36     end
37
38 end

```

### 9.2.3 assignmentStep.m

Listing 4: Helper function for `kmeans_alt(.)`.

```

1  %{
2  ECE 592 Project 1
3
4  Kudiyar (Cody) Orazymbetov
5  korazym@ncsu.edu
6
7  Nico Casale
8  ncasale@ncsu.edu
9
10 Helper function for kmeans_alt
11
12 Assigns each patch to a cluster,

```

```

13 Sums continuously the new centroid, to be averaged later
14 Counts the number of patches in each cluster
15 %}
16
17 function [Cnew, count, idx] = assignmentStep(Cn, idx, patches)
18
19 %{
20 assign each patch to one of the K centroids in Cn
21 based on minimum distance to a given centroid
22 %}
23
24 % todo: repmat on Cn to make more parallel?
25
26 % initialize new clusters
27 Cnew = zeros(size(Cn));
28 count = zeros(size(Cn,1),1);
29
30 for i = 1:length(patches)
31     % broken up step by step to find bottleneck
32     a1 = (patches(i) - Cn).^2; % least expensive
33     a2 = sum(a1,2); % mid-expensive
34     [~,newId] = min(a2); % most expensive
35
36     idx(i) = newId;
37     Cnew(newId,:) = Cnew(newId,:) + patches(i,:);
38     count(newId) = count(newId) + 1;
39 end
40
41 end

```

#### 9.2.4 updateStep.m

Listing 5: Helper function for `kmeans_alt(.)`.

```

1  %{
2  ECE 592 Project 1
3
4  Kudiyar (Cody) Orazymbetov
5  korazym@ncsu.edu
6
7  Nico Casale
8  ncasale@ncsu.edu
9
10 Helper function for kmeans_alt
11
12 Updates each cluster as the centroid of the patches
13 assigned to the cluster at hand
14 %}
15
16 function Cnew = updateStep(Cnew, count)
17
18 % average
19 for i = 1:length(Cnew)
20     % don't modify the cluster if no one belongs to it
21     if count(i) == 0
22         continue;

```



```
23         end
24         Cnew(i,:) = Cnew(i,+)/count(i);
25     end
26 end
```